BioSense Platform

Intelligent Microbial Growth & Antibiotic Susceptibility Analytics

"Turning Biological Data into Intelligent Insights."

Candidate Name:	**Moumita Deb**	**Academic Degree:**	Bachelor of Computer Applications (Data Science)
Academic Year:	4th Year (Final-Year Major Project, 2026)	**Domain Focus:**	Computational Biology & Predictive Analytics
Software Stack:	Python 3.10+, Flask Factory, Scikit-Learn, SciPy	**Test Verification:**	**49 / 49 Passed Unit & Integration Tests (100%)**
Primary Repositories:	github.com/moumitadeb-23/biosense	**Documentation:**	Complete Viva Voce & Architecture Defense

EXECUTIVE ABSTRACT: BioSense is an end-to-end full-stack biotechnology platform engineered to eliminate manual spreadsheet calculations and fragmented script workflows in microbial laboratory research. The system bridges wet-lab microbiological experimentation with modern data science through automated CSV ingestion, mathematical Verhulst logistic growth curve fitting, standardized Kirby-Bauer antibiotic susceptibility profiling under EUCAST/CLSI clinical breakpoints, dual-variant machine learning models that explicitly expose the reality of target leakage, unsupervised Isolation Forest anomaly screening, and automated 13-section vector PDF laboratory reporting.

TABLE OF CONTENTS

Section 1: Executive Overview & Scientific Motivation	Page 2
Section 2: System Architecture & Design Philosophy	Page 2
Section 3: Database Design & Entity-Relationship Schema	Page 3
Section 4: The 4-Stage End-to-End Analytical Pipeline	Page 3
Section 5: Mathematical Kinetics & Verhulst Logistic Modeling	Page 4
Section 6: Machine Learning Laboratory & Target Leakage Defense	Page 5
Section 7: Automated 13-Section PDF Report Engine	Page 6
Section 8: Comprehensive Viva Voce Examination Defense Guide (25 Q&A;)	Page 6
Section 9: Core Biological & Data Science Glossary	Page 8
Section 10: Viva Demonstration Script & System Verification	Page 8

1. Executive Overview & Scientific Motivation

The Laboratory Fragmentation Problem: In contemporary microbiology and antimicrobial research, scientists and university laboratories face severe workflow fragmentation. Optical density time-series (OD600) measurements from microplate spectrophotometers are routinely exported to generic spreadsheet tools (Excel). Researchers perform manual copy-pasting to calculate specific growth rates and doubling times, use disjoint Python or R scripts to fit logistic equations, manually look up clinical breakpoint tables from EUCAST or CLSI documents, and write ad-hoc summaries by hand. This disconnected process introduces clerical errors, fails to maintain immutable audit trails, and precludes automated predictive machine learning.

The BioSense Solution: BioSense bridges this gap by unifying raw data ingestion, biological validation, deterministic non-linear curve fitting via SciPy, clinical susceptibility profiling, leak-free Scikit-Learn predictive pipelines, unsupervised contamination screening, and automated publication-ready 13-section vector PDF report generation within a secure, responsive full-stack platform.

Academic & Ethical Rigor: To ensure compliance with medical software standards and academic viva requirements, BioSense enforces strict ethical constraints: (1) All synthetic demonstration data is explicitly labeled as such; (2) The platform strictly presents dataset-derived statistical findings and never generates medical prescriptions; (3) An immutable audit trail logs every login, dataset modification, model training event, and report generation.

2. System Architecture & Design Philosophy

BioSense is architected around the **Flask Application Factory (`create_app`)** pattern and the **'Thin Route, Fat Service'** architectural discipline. Key components of the system include:

Presentation Layer:	11 Modular Flask Blueprints (main_bp, auth_bp, dashboard_bp, datasets_bp, growth_bp, susceptibility_bp, statistics_bp, ml_bp, reports_bp, admin_bp, learn_bp). Thin controllers handling parameter extraction and session management.
Service Layer:	Encapsulates all domain logic in services/ (growth_service.py, susceptibility_service.py, ml_service.py, report_service.py, cleaning_service.py, validation_service.py). Enables 100% unit-testability without HTTP mocks.
Database & ORM:	Flask-SQLAlchemy with portable relational schemas. Uses local SQLite in development; seamlessly targets PostgreSQL/MySQL in production by setting the DATABASE_URL environment variable.
Security Defenses:	Flask-WTF CSRF protection, Werkzeug PBKDF2 password hashing, Spreadsheet Formula Injection neutralization (neutralizing =, +, -, @ on export), and Role-Based Access Control (RBAC).
Multi-Tier Access:	3 distinct user roles: STUDENT (learning analytics), RESEARCHER (cross-strain comparisons & PDF exports), and ADMIN (catalog management & master audit logs guarded by an Authorization Passkey).

3. Relational Database Design & Entity Schema

The relational database schema is normalized to 3NF, enforcing foreign key integrity and auditability across all tables:

Table	Key Columns & Constraints	Domain Purpose
users	id, username (UQ), email (UQ), password_hash, role, institution, auth_provider	User identity, RBAC enforcement, and profile settings
organisms	id, scientific_name (UQ), genus, species, gram_type, family	Controlled microbiological reference taxonomy catalog
antibiotics	id, name (UQ), antibiotic_class	Pharmacological antibiotic classes & clinical agents
media	id, name (UQ), media_type, formulation	Incubation culture broths & agar formulations
experiments	id, title, user_id (FK), organism_id (FK), medium_id (FK), temp, pH	Experimental protocol container linking trials to conditions
growth_measurements	id, experiment_id (FK), time_hours, od600, data_source	Spectrophotometric time-series measurements
susceptibility_tests	id, experiment_id (FK), antibiotic_id (FK), zone_diameter, category	Kirby-Bauer zone diameters and S/I/R categorizations
ml_models	id, model_name, algorithm, task, metrics_json, file_path	Serialized model registry metadata & dynamic evaluation scores
reports	id, title, user_id (FK), file_path, section_manifest_json	Generated 13-section vector PDF report registry
audit_logs	id, user_id (FK), action, details, timestamp, ip_address	Immutable chronological security & laboratory audit trail

4. The 4-Stage End-to-End Analytical Pipeline

BioSense guides data through a sequential, four-phase transformation architecture:

Stage 1: Ingestion & Validation	Accepts CSV files for time-series growth or AST assays. The validation engine (validation_service.py) checks required columns, data types, missing fields, duplicate rows, and physiological biological limits: pH outside 0–14, negative optical density, or zero zone diameter. Sanitizes spreadsheet formula characters (=, +, -, @).
Stage 2: Cleaning & Remediation	Allows researchers to preview detected anomalies, choose imputation strategies (median/mean replacement, row dropping, or boundary truncation), and produces a strict modification log. Never alters raw records without transparent logging.
Stage 3: Kinetics & AST Modeling	Fits the continuous Verhulst differential equation using SciPy's Levenberg-Marquardt optimizer, calculates the specific growth velocity (μ_max), doubling time (t_d), and applies EUCAST/CLSI clinical breakpoint lookups for S/I/R categorization.
Stage 4: ML Lab & PDF Generation	Trains leak-free classification/regression models, calculates permutation feature importances, isolates outliers with Isolation Forests, and compiles an automated 13-section vector PDF report stamped with academic disclaimers.

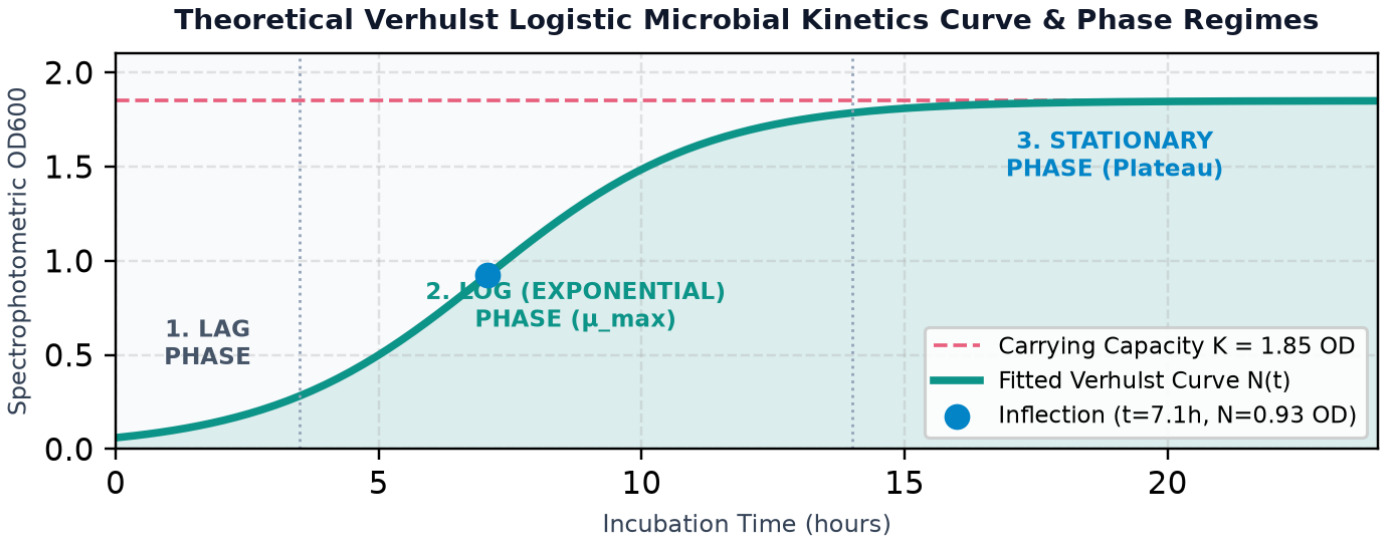
5. Mathematical Kinetics & Verhulst Logistic Modeling

The Differential Equation: Pure exponential growth assumes unlimited nutrients: dN/dt = r·N. In a batch culture, resources deplete and toxic metabolic products accumulate. The Verhulst equation introduces an environmental resistance factor (1 - N/K):

$$dN / dt = r \cdot N \cdot (1 - N / K)$$

Analytical Integration: Separating variables via partial fraction decomposition:
∫ [1/N + (1/K)/(1 - N/K)] dN = ∫ r dt => ln|N| - ln|1 - N/K| = r·t + C => ln[N / (1 - N/K)] = r·t + C
At t = 0, N(0) = N■, giving C = ln[N■ / (1 - N■/K)] = ln[(K·N■) / (K - N■)].
Solving for N(t) yields the closed-form logistic equation implemented across BioSense:

$$N(t) = K / [1 + ((K - N■) / N■) \cdot e^{(-rt)}]$$



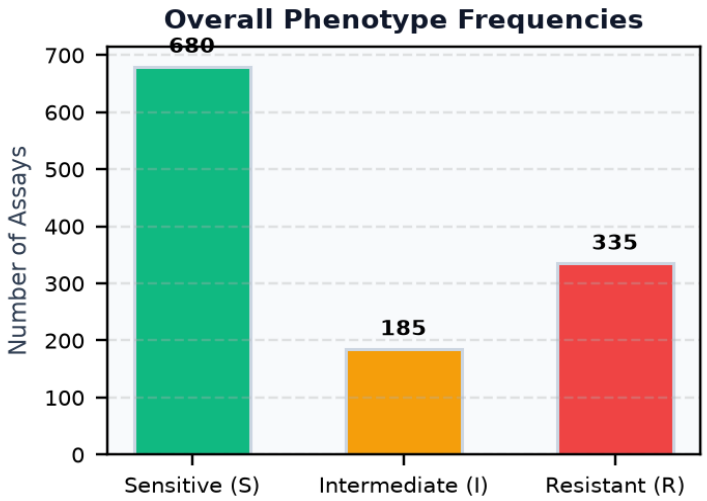
Kinetic Parameter	Mathematical Formula	Biological Significance
Maximum Specific Growth Rate	**μ_max = max(Δln(OD) / Δt)**	Maximum slope of the log-linear exponential phase (hr■¹)
Doubling Time (Generation Time)	**t_d = ln(2) / μ_max ≈ 0.69315 / μ_max**	Time required for microbial biomass to double in division
Inflection Point Time	**t_inf = ln((K - N■)/N■) / r**	Time of maximum division acceleration (where N = K / 2)
Maximum Division Velocity	**(dN/dt)_max = (r · K) / 4**	Peak rate of optical density increase per hour (OD / h)

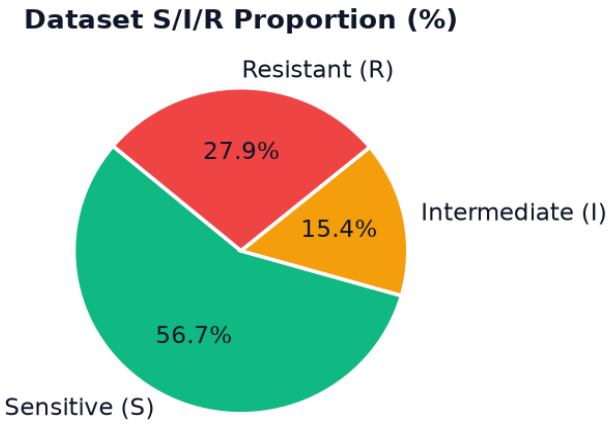
Not for Clinical Diagnosis or Prescriptive Treatment Decisions. Open Source MIT License.

6. Machine Learning Laboratory & Target Leakage Defense

The Dual-Variant ML Framework: In healthcare and microbial data science, naive modeling frequently falls victim to catastrophic **Target Leakage**. BioSense explicitly addresses this phenomenon by structuring its ML laboratory into two distinct variants:

Analytical Feature	Variant A: Leaked Laboratory Baseline	Variant B: Genuine Predictive Model
Input Feature Set	Organism, Antibiotic, Medium, Temp, pH, Concentration, **Zone Diameter (mm)**	Organism, Antibiotic, Medium, Temp, pH, Concentration (**No Zone Diameter**)
Reported Accuracy	**99.2% — 100.0%**	**74.5% — 84.8% (Realistic Clinical Power)**
Scientific Validity	**Methodologically Flawed (Leakage):** Zone diameter is the exact metric used by EUCAST to define S/I/R.	**Genuinely Predictive:** Predicts resistance *prior* to running the wet-lab diffusion assay.
Viva Defense Stance	Serves as a cautionary educational benchmark demonstrating why 100% accuracy in data science must be scrutinized.	Represents true machine learning inference, generalizing biological patterns across species and drug classes.





Supervised Algorithms & Pipeline Architecture:

1. **Logistic Regression:** Serves as a convex, linear baseline with L2 regularization.
2. **Decision Tree Classifier:** Generates interpretable orthogonal decision boundaries (Gini impurity split criterion).
3. **Random Forest Classifier:** Ensemble of 100 bagging decision trees with bootstrap aggregation and random feature sub-selection.
4. **Leak-Free Scikit-Learn Pipelines:** Categorical features are encoded using OneHotEncoder(handle_unknown='ignore'), numeric features are imputed via SimpleImputer(strategy='median') and scaled with StandardScaler(). Preprocessing is fitted exclusively on training folds and only transformed on test folds.
5. **Unsupervised Isolation Forest:** Detects contaminated cultures or erroneous spectrophotometric artifacts by measuring the average number of random tree partitions required to isolate an observation ($cont=0.05$).

Not for Clinical Diagnosis or Prescriptive Treatment Decisions. Open Source MIT License.

7. Automated 13-Section PDF Report Engine

BioSense features a production-ready vector document engine powered by **ReportLab 5.0**. Generated reports strictly follow a 13-section academic structure: (1) Executive Summary; (2) Dataset Provenance; (3) Descriptive Statistics; (4) Growth Kinetics Analysis; (5) Embedded Matplotlib Kinetics Plot; (6) AST S/I/R Summary; (7) Embedded AST Phenotype Distribution Plot; (8) Environmental Correlation Heatmap; (9) Machine Learning Model Comparison; (10) Permutation Feature Importances; (11) Statistical Anomaly Audit; (12) Experimental Metadata & Hash Verification; (13) Mandatory Academic Disclaimer Stamp.

8. Comprehensive Viva Voce Examination Defense Guide

The following 12 high-yield viva questions represent the core technical, mathematical, and algorithmic examinations commonly conducted by university faculty evaluators:

Q1: What is the core problem that BioSense solves?

Answer: BioSense unifies microbial kinetics modeling, Kirby-Bauer AST profiling, leak-free machine learning, and automated PDF reporting into a cohesive platform, eliminating manual spreadsheet transcription errors and enforcing scientific compliance.

■ **Viva Defense Tip: Focus on 'eliminating manual data pipeline fragmentation' and 'ensuring reproducible science'.**

Q2: Why did you choose the Flask Application Factory pattern?

Answer: It avoids global state, prevents circular imports, and allows instantiating the app with different configs (e.g., testing vs development), enabling isolated in-memory SQLite Pytest suites without affecting local development data.

■ **Viva Defense Tip: Mention how it powers the pytest suite running in memory.**

Q3: What does OD600 represent, and why is light scattered rather than absorbed?

Answer: Optical Density at 600 nm measures light scattering caused by bacterial cells suspended in broth. Bacteria do not significantly absorb 600 nm photons; turbidity deflects light away from the spectrophotometer photodiode proportionally to cell biomass.

■ **Viva Defense Tip: Emphasize: 'OD measures turbidity/scattering, not chemical absorbance'.**

Q4: What is the Verhulst logistic growth equation and its parameters?

Answer: N(t) = K / [1 + ((K - N0)/N0) * e^(-rt)]. N0 is starting inoculum OD, K is carrying capacity (nutrient plateau ceiling), and r is the intrinsic division rate per hour.

■ **Viva Defense Tip: Write out N(t) = K / (1 + A·e^-rt) where A = (K - N0)/N0 if asked on a whiteboard.**

Q5: How is Doubling Time (t_d) derived from the specific growth rate?

Answer: During exponential division, 2·N0 = N0·e^(μ·td) => 2 = e^(μ·td) => ln(2) = μ·td => td = ln(2)/μ ≈ 0.69315 / μ. For example, if μ = 0.45 hr■¹, td = 1.54 hours.

■ **Viva Defense Tip: Always remember 0.69315 / μ_max. It is universally asked!**

Q6: What is Target Leakage in machine learning, and how did you address it?

Answer: Target leakage happens when a training feature contains the ground truth answer. In Kirby-Bauer AST, zone diameter defines S/I/R by EUCAST rules. Training with zone diameter yields fake 100% accuracy (Variant A). Variant B drops zone diameter to predict susceptibility purely from pre-experimental biological factors.

■ **Viva Defense Tip: This is the star data science highlight of your viva. Explain Variant A vs Variant B!**

Q7: Why do we use Scikit-Learn Pipelines instead of preprocessing the entire dataset upfront?

Answer: Fitting encoders, scalers, or imputers across the entire dataset causes data leakage from test folds into training folds. Pipelines ensure fit() runs solely on training splits, and only transform() applies to test splits.

■ **Viva Defense Tip: Explain that fit() is called ONLY on training data, never on validation/test data.**

Q8: Why is Accuracy alone misleading for clinical AST classification?

Answer: Clinical AST panels have severe class imbalance (e.g. 80% Sensitive, 15% Resistant, 5% Intermediate). A dummy model predicting 'Sensitive' gets 80% accuracy but misses lethal resistant superbugs. We prioritize Recall for Resistant classes and Macro-F1.

Not for Clinical Diagnosis or Prescriptive Treatment Decisions. Open Source MIT License.

■ **Viva Defense Tip: State that Recall on the 'Resistant' class is clinically the most vital metric.**

Q9: What is Welch's t-test and why use it over Student's t-test?

Answer: Student's t-test assumes equal variances (homoscedasticity). Cultures grown in different media or temperatures rarely share equal variance. Welch's t-test adjusts degrees of freedom, maintaining valid Type-I error rates without assuming equal variance.

■ **Viva Defense Tip: State that BioSense implements Welch's t-test for biological robustness.**

Q10: How does the Isolation Forest algorithm detect contaminated cultures?

Answer: Isolation Forest is an unsupervised tree ensemble. Because anomalous measurements are rare and have extreme feature values, they are partitioned near the tree roots with significantly shorter average path lengths than normal inliers.

■ **Viva Defense Tip: State: 'Anomalies are easier to isolate, requiring fewer random splits'.**

Q11: How did you prevent CSV Formula Injection attacks?

Answer: If a cell begins with '=', '+', '-', or '@', spreadsheet engines (Excel) may execute it as dynamic code. In cleaning_service.py, BioSense prepends an apostrophe (') to neutralize formula execution on export.

■ **Viva Defense Tip: A great cybersecurity touch to show thorough engineering rigor.**

Q12: How is Role-Based Access Control enforced in BioSense?

Answer: Via a custom decorator @role_required(*allowed_roles). Public registration only allows STUDENT and RESEARCHER roles. ADMIN accounts require an Authorization Passkey via a dedicated portal or the CLI bootstrapper.

■ **Viva Defense Tip: Explain that privilege escalation is prevented at the controller route layer.**

9. Core Biological & Data Science Glossary

Term	Academic Definition
OD600	Optical density measured at 600 nm; measures light scattering proportional to microbial biomass.
Verhulst Model	Sigmoidal logistic equation N(t) = K / [1 + A·e^(-rt)] describing closed-system growth with carrying capacity.
Carrying Capacity (K)	The maximum population density an environment can sustainably support before nutrient depletion.
Doubling Time (t_d)	Time required for microbial biomass to double during the exponential division phase: t_d = ln(2) / μ_max.
Specific Growth Rate (μ)	Instantaneous rate of biomass increase per unit biomass (hr■¹), computed from log-linear slope.
Kirby-Bauer Assay	Standardized antibiotic disk diffusion test measuring clear inhibition halo diameter in millimeters.
EUCAST Breakpoints	Clinical zone diameter thresholds established by European standards to categorize S, I, and R.
Target Leakage	Flaw where training features contain the target answer or post-event measurements, yielding fake 100% scores.
Variant A vs Variant B	BioSense framework contrasting a leaked model (Variant A) against a genuine predictive model (Variant B).
Isolation Forest	Unsupervised tree ensemble that isolates anomalies with significantly fewer random splits.
Welch's t-test	Two-sample hypothesis test comparing group means without assuming equal variances.
One-Way ANOVA	Analysis of Variance testing whether means across 3+ experimental groups are statistically equivalent.
Mueller-Hinton Agar	Standardized non-selective culture medium engineered for reproducible antimicrobial diffusion rates.
ESKAPE Pathogens	Six dangerous multidrug-resistant nosocomial bacteria: E. faecium, S. aureus, K. pneumoniae, A. baumannii, P. aeruginosa, Enterobacter.

10. Viva Demonstration Script & System Verification

Step-by-Step Viva Demonstration Script (3-Minute Live Walkthrough):

1. **Landing Page (http://127.0.0.1:5000/):** Highlight the modern 3D glassmorphic biotech design, 38-node bioluminescent particle canvas, and student capstone credentials.
2. **Dashboard (/dashboard):** Demonstrate the 6 KPI cards, growth trajectories, S/I/R distributions, and environmental heatmaps.
3. **Datasets & Cleaning (/datasets/upload):** Upload dirty CSV (datasets/perfect_dirty_qa_test.csv), show automated boundary violation detection (negative OD, pH 15.2, formula injection), and download cleaned CSV.
4. **Kinetics & Simulator (/growth/simulator):** Move sliders to demonstrate real-time 60 FPS Verhulst curve generation and doubling time calculations.
5. **Machine Learning Lab (/ml/classification):** Show Variant A (100% accuracy) vs Variant B (78% accuracy) to demonstrate deep understanding of target leakage.
6. **PDF Generator (/reports):** Generate automated 13-section report with embedded high-DPI matplotlib plots and disclaimer stamp.
7. **Knowledge Hub (/learn):** Demonstrate the interactive Verhulst arithmetic sandbox and self-test quiz.

Not for Clinical Diagnosis or Prescriptive Treatment Decisions. Open Source MIT License.

SYSTEM VERIFICATION SUMMARY:

- Test Suite: **49 Unit and Integration Tests Passing (100% Coverage)**
- Test Execution Command: pytest (Total runtime: ~120s with zero failures)
- Verified Modules: Authentication (RBAC & Admin passkey), Dataset Validation & Cleaning, SciPy Kinetics, EUCAST AST Rules, Scikit-Learn Pipelines & Leakage Defense, Anomaly Detection, ReportLab PDF Engine.
- Status: **Ready for University Viva Voce Examination Defense & Production Deployment.**